

POLICY HYSTERESIS IN CONTACT TASKS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study contact-rich robot policy learning. The motivating bet is: Expose and exploit path-dependent contact states ignored by memoryless policies. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For contact-rich robot policy learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2025a; 2021; 2025b; 2024; 2025c; 2023). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.500	0.500
observed_only_wam	hidden_mode_shift	0.168	0.833
augmentation_as_data_wam	nominal	0.625	0.375
augmentation_as_data_wam	hidden_mode_shift	0.263	0.738
scalar_uncertainty_wam	nominal	0.645	0.355
scalar_uncertainty_wam	hidden_mode_shift	0.340	0.660
proposed_branch_mechanism	nominal	0.802	0.198
proposed_branch_mechanism	hidden_mode_shift	0.547	0.453

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Policy Hysteresis in Contact Tasks argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Vision-guided mpc for robotic path following using learned memory-augmented model, 2021. <https://doi.org/10.3389/frobt.2021.688275>.
- Vision-force-fused curriculum learning for robotic contact-rich assembly tasks, 2023. <https://doi.org/10.3389/fnbot.2023.1280773>.
- Modeling of slip rate-dependent traversability for path planning of wheeled mobile robot in sandy terrain, 2024. <https://doi.org/10.3389/frobt.2024.1320261>.
- Force-dependent variable impedance controller for contact-rich tasks under reference trajectory uncertainty, 2025a. <https://doi.org/10.1109/lcsys.2025.3597334>.
- Polytouch: A robust multi-modal tactile sensor for contact-rich manipulation using tactile-diffusion policies, 2025b. <http://arxiv.org/abs/2504.19341v1>.
- Share-rl: Structured, interactive reinforcement learning for contact-rich industrial assembly tasks, 2025c. <https://doi.org/10.48550/arxiv.2509.13949>.